

```
In [1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load in

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# import other needed packages and functions
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.preprocessing import MinMaxScaler
from sklearn.feature_selection import SelectKBest
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier, AdaBoostClassifier, GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import VotingClassifier
import lightgbm
from sklearn.model_selection import cross_val_score
import itertools

# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# Any results you write to the current directory are saved as output.

/kaggle/input/titanic/train.csv
/kaggle/input/titanic/gender_submission.csv
/kaggle/input/titanic/test.csv
```

```
In [2]: # read in csv files and make dfs
train_df = pd.read_csv("/kaggle/input/titanic/train.csv")
test_df = pd.read_csv("/kaggle/input/titanic/test.csv")

# make copy of test df for submission
submission = test_df.copy()

# combine train and test dfs into 1 df of all data
all_df = pd.concat([train_df, test_df], sort=False)
```

```
In [3]: # Define function to inspect data frames. Prints first few lines, determines s
        # ize/shape of data frame,
        # shows descriptive statistics, shows data types, shows missing or incomplete
        # data, check for duplicate data.

def inspect_df(df):
    print('Header:')
    print('{}'.format(df.head()))
    print()
    print('Shape: {}'.format(df.shape))
    print()
    print('Statistics:')
    print('{}'.format(df.describe()))
    print()
    print('Info:')
    print('{}'.format(df.info()))

# use inspect_df on all_df

inspect_df(all_df)
```

Header:

	PassengerId	Survived	Pclass	\
0	1	0.0	3	
1	2	1.0	1	
2	3	1.0	3	
3	4	1.0	1	
4	5	0.0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

Shape: (1309, 12)

Statistics:

	PassengerId	Survived	Pclass	Age	SibSp	\
count	1309.000000	891.000000	1309.000000	1046.000000	1309.000000	
mean	655.000000	0.383838	2.294882	29.881138	0.498854	
std	378.020061	0.486592	0.837836	14.413493	1.041658	
min	1.000000	0.000000	1.000000	0.170000	0.000000	
25%	328.000000	0.000000	2.000000	21.000000	0.000000	
50%	655.000000	0.000000	3.000000	28.000000	0.000000	
75%	982.000000	1.000000	3.000000	39.000000	1.000000	
max	1309.000000	1.000000	3.000000	80.000000	8.000000	

	Parch	Fare
count	1309.000000	1308.000000
mean	0.385027	33.295479
std	0.865560	51.758668
min	0.000000	0.000000
25%	0.000000	7.895800
50%	0.000000	14.454200
75%	0.000000	31.275000
max	9.000000	512.329200

Info:

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1309 entries, 0 to 417
Data columns (total 12 columns):
PassengerId    1309 non-null int64
Survived       891 non-null float64
Pclass         1309 non-null int64
Name           1309 non-null object
Sex            1309 non-null object
Age           1046 non-null float64
SibSp          1309 non-null int64
Parch          1309 non-null int64
```

```

Ticket      1309 non-null object
Fare        1308 non-null float64
Cabin       295 non-null object
Embarked    1307 non-null object
dtypes: float64(3), int64(4), object(5)
memory usage: 132.9+ KB
None

```

In [4]: *# Look at proportions of passengers by Pclass*

```
all_df.Pclass.value_counts(normalize=True, sort=False)
```

```

Out[4]: 1    0.246753
        2    0.211612
        3    0.541635
        Name: Pclass, dtype: float64

```

In [5]: *# Look at proportions of passengers by Sex*

```
all_df.Sex.value_counts(normalize=True)
```

```

Out[5]: male      0.644003
        female    0.355997
        Name: Sex, dtype: float64

```

In [6]: *# inspect null values for Embarked*

```
all_df[all_df.Embarked.isnull()]
```

Out[6]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
61	62	1.0	1	Icard, Miss. Amelie	female	38.0	0	0	113572	80.0	B28
829	830	1.0	1	Stone, Mrs. George Nelson (Martha Evelyn)	female	62.0	0	0	113572	80.0	B28

Looked up Mrs. Stone and Miss Icard online, they boarded in Southampton.

In [7]: *# fill missing values for Embarked with information found online*

```
all_df.loc[[61, 829], ['Embarked']] = 'S'
```

In [8]: *# plot histogram of Ages*


```
plt.hist(data = all_df, x = 'Age', bins = 40);
```

/opt/conda/lib/python3.6/site-packages/numpy/lib/histograms.py:824: RuntimeWarning: invalid value encountered in greater_equal

```
    keep = (tmp_a >= first_edge)
```


/opt/conda/lib/python3.6/site-packages/numpy/lib/histograms.py:825: RuntimeWarning: invalid value encountered in less_equal

```
    keep &= (tmp_a <= last_edge)
```




In [9]: *# plot histogram of Fare*

```
plt.hist(data = all_df, x = 'Fare', bins = 40);
```




```
In [10]: # plot histogram of siblings and spouses
```

```
plt.hist(data = all_df, x = 'SibSp', bins = 8);
```



```
In [11]: # plot histogram of parents and children
```

```
plt.hist(data = all_df, x = 'Parch', bins = 9);
```



In [12]: *# inspect columns and missing values again*

```
all_df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1309 entries, 0 to 417
Data columns (total 12 columns):
PassengerId    1309 non-null int64
Survived       891 non-null float64
Pclass         1309 non-null int64
Name           1309 non-null object
Sex            1309 non-null object
Age            1046 non-null float64
SibSp          1309 non-null int64
Parch          1309 non-null int64
Ticket         1309 non-null object
Fare           1308 non-null float64
Cabin          295 non-null object
Embarked       1309 non-null object
dtypes: float64(3), int64(4), object(5)
memory usage: 172.9+ KB
```

In [13]: *# find passenger with missing fare data*

```
all_df[all_df.Fare.isnull()]
```

Out[13]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Er
152	1044	NaN	3	Storey, Mr. Thomas	male	60.5	0	0	3701	NaN	NaN	

In [14]: *# get average fare of each Pclass*

```
all_df.Fare.groupby(all_df.Pclass).mean()
```

Out[14]: Pclass

```
1    87.508992
2    21.179196
3    13.302889
Name: Fare, dtype: float64
```

In [15]: *#fill nan fare value with rounded mean for class 3*

```
all_df.loc[152, ['Fare']] = 13
```

In [16]: *# fill missing values for ages with the mean age value for each passengers Pclass and Sex*

```
all_df.Age = all_df.Age.groupby([all_df.Pclass, all_df.Sex]).transform(lambda x: x.fillna(x.mean()))
```

```
In [17]: # Extract titles from names and make new Title column, then get survival rate of each title
```

```
all_df['Title'] = all_df.Name.str.extract(' ([A-Za-z]+)\.', expand=False)  
all_df.Survived.groupby(all_df.Title).mean()
```

```
Out[17]: Title  
Capt      0.000000  
Col        0.500000  
Countess   1.000000  
Don        0.000000  
Dona       NaN  
Dr         0.428571  
Jonkheer   0.000000  
Lady       1.000000  
Major      0.500000  
Master     0.575000  
Miss       0.697802  
Mlle       1.000000  
Mme        1.000000  
Mr         0.156673  
Mrs        0.792000  
Ms         1.000000  
Rev        0.000000  
Sir        1.000000  
Name: Survived, dtype: float64
```

```
In [18]: # get value counts for title occurrences
```

```
all_df.Title.value_counts()
```

```
Out[18]: Mr      757  
Miss    260  
Mrs     197  
Master   61  
Rev       8  
Dr        8  
Col       4  
Major     2  
Mlle      2  
Ms        2  
Sir       1  
Don       1  
Lady      1  
Mme       1  
Dona      1  
Jonkheer  1  
Capt     1  
Countess  1  
Name: Title, dtype: int64
```



```
In [19]: # Replace uncommon titles with more common values and view new occurrences

all_df.Title = all_df.Title.replace(['Capt', 'Col', 'Dr', 'Major', 'Rev', 'Don', 'Sir', 'Jonkheer'], 'Mr')
all_df.Title = all_df.Title.replace(['Ms', 'Mlle'], 'Miss')
all_df.Title = all_df.Title.replace(['Mme', 'Lady', 'Countess', 'Dona'], 'Mrs')
all_df.Title.value_counts()
```

```
Out[19]: Mr      783
Miss    264
Mrs     201
Master   61
Name: Title, dtype: int64
```

```
In [20]: # combine sibsp and parch to one family column
all_df['Fam'] = all_df.SibSp + all_df.Parch

# make ticket frequency column for number of occurrences of ticket number
all_df['Ticket_Frequency'] = all_df.groupby('Ticket')['Ticket'].transform('count')

# make column for solo vs travel with family
all_df.loc[all_df['Fam'] == 0, 'Solo'] = 1
all_df.loc[all_df['Ticket_Frequency'] == 1, 'Solo'] = 1
all_df.Solo = all_df.Solo.fillna(0)

# bin fare column to 9 quantiles and encode as ordinal
all_df['Fare'] = pd.qcut(all_df.Fare, q=9, labels=np.arange(1,10))

# bin age column to 10 quantiles and encode as ordinal
all_df['Age'] = pd.qcut(all_df.Age, q=10, labels=np.arange(1,11))

# one-hot encode sex column and capitalize sex columns for consistency
all_df = pd.concat([all_df, pd.get_dummies(all_df.Sex)], axis=1)
all_df.rename(columns={'male':'Male', 'female':'Female'}, inplace=True)

# one-hot encode embarked column
all_df = pd.concat([all_df, pd.get_dummies(all_df.Embarked, prefix='Embarked')], axis=1)

# one-hot encode title column
all_df = pd.concat([all_df, pd.get_dummies(all_df.Title)], axis=1)

# drop unwanted columns (name, sex, cabin, embarked and title have been replaced with one hot encoding, ticket replaced with ticket frequency,
# cabin has too many missing values, sibsp and parch replaced with fam and solo columns)
all_df = all_df.drop(columns=['Name', 'Sex', 'Ticket', 'Cabin', 'SibSp', 'Parch', 'Embarked', 'Title'])

# inspect columns and number of values for resulting df
all_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1309 entries, 0 to 417
Data columns (total 17 columns):
PassengerId      1309 non-null int64
Survived          891 non-null float64
Pclass           1309 non-null int64
Age              1309 non-null category
Fare             1309 non-null category
Fam              1309 non-null int64
Ticket_Frequency 1309 non-null int64
Solo             1309 non-null float64
Female           1309 non-null uint8
Male             1309 non-null uint8
Embarked_C       1309 non-null uint8
Embarked_Q       1309 non-null uint8
Embarked_S       1309 non-null uint8
Master           1309 non-null uint8
Miss             1309 non-null uint8
Mr               1309 non-null uint8
Mrs              1309 non-null uint8
dtypes: category(2), float64(2), int64(4), uint8(9)
memory usage: 86.4 KB

```

In [21]: *# inspect survival rates for number of family members onboard*

```
all_df.Survived.groupby(all_df.Fam).mean()
```

Out[21]: Fam

```

0      0.303538
1      0.552795
2      0.578431
3      0.724138
4      0.200000
5      0.136364
6      0.333333
7      0.000000
10     0.000000
Name: Survived, dtype: float64

```

In [22]: *# bin fam column values and get value counts*

```

all_df.Fam = pd.cut(all_df.Fam, bins=[0, 1, 4, 7, 11], include_lowest=True, right=False, labels=[1, 2, 3, 4])
all_df.Fam.value_counts()

```

Out[22]: 1 790

```

2      437
3       63
4       19
Name: Fam, dtype: int64

```

In [23]: *# inspect survival rate for number of people traveling in group*

```
all_df.Survived.groupby(all_df.Ticket_Frequency).mean()
```

Out[23]: Ticket_Frequency

```
1    0.270270
2    0.513812
3    0.653465
4    0.727273
5    0.333333
6    0.210526
7    0.208333
8    0.384615
11   0.000000
Name: Survived, dtype: float64
```

In [24]: *# bin ticket frequency column values and get value counts*

```
all_df.Ticket_Frequency = pd.cut(all_df.Ticket_Frequency, bins=[0, 2, 5, 9, 12], right=False, labels=[1, 2, 3, 4])
all_df.Ticket_Frequency.value_counts()
```

Out[24]:

```
1    713
2    475
3    110
4     11
Name: Ticket_Frequency, dtype: int64
```

In [25]: *# make list of all columns and view it*

```
cols = list(all_df)
cols
```

Out[25]:

```
['PassengerId',
 'Survived',
 'Pclass',
 'Age',
 'Fare',
 'Fam',
 'Ticket_Frequency',
 'Solo',
 'Female',
 'Male',
 'Embarked_C',
 'Embarked_Q',
 'Embarked_S',
 'Master',
 'Miss',
 'Mr',
 'Mrs']
```

In [26]: *# use min max scaler to scale all feature columns to range 0-1*

```
scaler = MinMaxScaler()
all_df[cols] = scaler.fit_transform(all_df[cols])
```

```
In [27]: # make array of survived labels (training), drop survived column from all_df,
         # split all_df into features (training) and test_df, and make array of feature
         # column names

labels = all_df.loc[:890, 'Survived']
all_df = all_df.drop(columns = 'Survived')
features = all_df.iloc[:891]
test_df = all_df.iloc[891:]
feat_names = features.columns.values
```

```
In [28]: # use SelectKBest to narrow down to top features and use result to transform t
         # rain and test features dfs

k = SelectKBest(k=11)
k.fit(features, labels)
k_scores = (k.scores_)
features = k.transform(features)
test_df = k.transform(test_df)
```

```
In [29]: # make df to show scores for all features and print

feat_scores = pd.DataFrame()
feat_scores['Feature'] = feat_names
feat_scores['Score'] = k_scores
feat_scores
```

Out[29]:

	Feature	Score
0	PassengerId	0.022285
1	Pclass	115.031272
2	Age	0.118749
3	Fare	108.327062
4	Fam	6.313788
5	Ticket_Frequency	17.149390
6	Solo	47.368609
7	Female	372.405724
8	Male	372.405724
9	Embarked_C	25.895987
10	Embarked_Q	0.011846
11	Embarked_S	20.374460
12	Master	6.503635
13	Miss	112.860827
14	Mr	414.442624
15	Mrs	122.387505

```
In [30]: # split training data into train and test subsets for validation

features_train, features_test, labels_train, labels_test = train_test_split(features, labels, test_size=0.2, random_state=3)
```

```
In [31]: # setup base classifiers

gb = GradientBoostingClassifier()
rf = RandomForestClassifier()
et = ExtraTreesClassifier()
ab = AdaBoostClassifier()
dt = DecisionTreeClassifier()
lr = LogisticRegression(solver='liblinear')
kn = KNeighborsClassifier()
svc = SVC(gamma='auto', probability=True)
gnb = GaussianNB()
```

In [32]: *# get train base classifiers and get initial validation scores*

```
gb.fit(features_train, labels_train)
print('GB Score:', gb.score(features_test, labels_test))
rf.fit(features_train, labels_train)
print('RF Score:', rf.score(features_test, labels_test))
et.fit(features_train, labels_train)
print('ET Score:', et.score(features_test, labels_test))
ab.fit(features_train, labels_train)
print('AB Score:', ab.score(features_test, labels_test))
dt.fit(features_train, labels_train)
print('DT Score:', dt.score(features_test, labels_test))
lr.fit(features_train, labels_train)
print('LR Score:', lr.score(features_test, labels_test))
kn.fit(features_train, labels_train)
print('KN Score:', kn.score(features_test, labels_test))
svc.fit(features_train, labels_train)
print('SVC Score:', svc.score(features_test, labels_test))
gnb.fit(features_train, labels_train)
print('GNB Score:', gnb.score(features_test, labels_test))
```

GB Score: 0.8156424581005587

RF Score: 0.8212290502793296

ET Score: 0.8156424581005587

AB Score: 0.8044692737430168

DT Score: 0.8156424581005587

LR Score: 0.8212290502793296

KN Score: 0.8044692737430168

/opt/conda/lib/python3.6/site-packages/sklearn/ensemble/forest.py:245: Future Warning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.

"10 in version 0.20 to 100 in 0.22.", FutureWarning)

/opt/conda/lib/python3.6/site-packages/sklearn/ensemble/forest.py:245: Future Warning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.

"10 in version 0.20 to 100 in 0.22.", FutureWarning)

SVC Score: 0.8044692737430168

GNB Score: 0.7821229050279329

In [33]: *# use gridsearch to tune base classifier hyperparameters*

```
alg = gb
params = {'n_estimators': (10, 25, 50, 100), 'learning_rate': (0.01, 0.1, 0.5, 1, 5, 10)}
clf = GridSearchCV(alg, params, cv = 5, scoring = 'accuracy', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
print("Best Score:", clf.best_score_)
gb = clf.best_estimator_
```

Best Parameters: {'learning_rate': 0.01, 'n_estimators': 100}
Best Score: 0.8286516853932584

/opt/conda/lib/python3.6/site-packages/sklearn/model_selection/_search.py:81
4: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
DeprecationWarning)

In [34]: *# use gridsearch to tune base classifier hyperparameters*

```
alg = rf
params = {'n_estimators': (10, 25, 50, 100), 'min_samples_split': (2, 3, 4, 5, 10), 'min_samples_leaf': (1, 2, 3, 4, 5)}
clf = GridSearchCV(alg, params, cv = 5, scoring = 'accuracy', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
print("Best Score:", clf.best_score_)
rf = clf.best_estimator_
```

Best Parameters: {'min_samples_leaf': 5, 'min_samples_split': 3, 'n_estimators': 25}
Best Score: 0.8300561797752809

In [35]: *# use gridsearch to tune base classifier hyperparameters*

```
alg = et
params = {'n_estimators': (10, 25, 50, 100), 'min_samples_split': (2, 3, 4, 5, 10), 'min_samples_leaf': (1, 2, 3, 4, 5)}
clf = GridSearchCV(alg, params, cv = 5, scoring = 'accuracy', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
print("Best Score:", clf.best_score_)
et = clf.best_estimator_
```

Best Parameters: {'min_samples_leaf': 4, 'min_samples_split': 4, 'n_estimators': 10}
Best Score: 0.8370786516853933


```
In [36]: # use gridsearch to tune base classifier hyperparameters

alg = ab
params = {'n_estimators': (10, 25, 50, 100), 'learning_rate': (0.01, 0.1, 0.5, 1, 5, 10)}
clf = GridSearchCV(alg, params, cv = 5, scoring = 'accuracy', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
print("Best Score:", clf.best_score_)
ab = clf.best_estimator_
```

Best Parameters: {'learning_rate': 1, 'n_estimators': 10}
Best Score: 0.8146067415730337

```
In [37]: # use gridsearch to tune base classifier hyperparameters

alg = dt
params = {'min_samples_split': (2, 3, 4, 5, 10), 'min_samples_leaf': (1, 2, 3, 4, 5)}
clf = GridSearchCV(alg, params, cv = 5, scoring = 'accuracy', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
print("Best Score:", clf.best_score_)
dt = clf.best_estimator_
```

Best Parameters: {'min_samples_leaf': 1, 'min_samples_split': 4}
Best Score: 0.8174157303370787

```
In [38]: # use gridsearch to tune base classifier hyperparameters

alg = lr
params = {'penalty': ('l1', 'l2'), 'C': (0.01, 0.1, 0.5, 1, 5, 10), 'max_iter': (100, 500)}
clf = GridSearchCV(alg, params, cv = 5, scoring = 'accuracy', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
print("Best Score:", clf.best_score_)
lr = clf.best_estimator_
```

Best Parameters: {'C': 10, 'max_iter': 100, 'penalty': 'l1'}
Best Score: 0.8188202247191011

```
In [39]: # use gridsearch to tune base classifier hyperparameters

alg = kn
params = {'n_neighbors': (2, 3, 4, 5, 10, 20)}
clf = GridSearchCV(alg, params, cv = 5, scoring = 'accuracy', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
print("Best Score:", clf.best_score_)
kn = clf.best_estimator_
```

Best Parameters: {'n_neighbors': 10}
Best Score: 0.8117977528089888

```
In [40]: # use gridsearch to tune base classifier hyperparameters

alg = svc
params = {'C': (0.01, 0.1, 0.5, 1, 5, 10), 'kernel': ('linear', 'poly', 'rbf',
'sigmoid')}
clf = GridSearchCV(alg, params, cv = 5, scoring = 'accuracy', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
print("Best Score:", clf.best_score_)
svc = clf.best_estimator_
```

Best Parameters: {'C': 10, 'kernel': 'poly'}
Best Score: 0.824438202247191

```
In [41]: # retrain base classifiers and get validation scores

gb.fit(features_train, labels_train)
print('GB Score:', gb.score(features_test, labels_test))
rf.fit(features_train, labels_train)
print('RF Score:', rf.score(features_test, labels_test))
et.fit(features_train, labels_train)
print('ET Score:', et.score(features_test, labels_test))
ab.fit(features_train, labels_train)
print('AB Score:', ab.score(features_test, labels_test))
dt.fit(features_train, labels_train)
print('DT Score:', dt.score(features_test, labels_test))
lr.fit(features_train, labels_train)
print('LR Score:', lr.score(features_test, labels_test))
kn.fit(features_train, labels_train)
print('KN Score:', kn.score(features_test, labels_test))
svc.fit(features_train, labels_train)
print('SVC Score:', svc.score(features_test, labels_test))
gnb.fit(features_train, labels_train)
print('GNB Score:', gnb.score(features_test, labels_test))
```

GB Score: 0.8379888268156425
RF Score: 0.8379888268156425
ET Score: 0.8212290502793296
AB Score: 0.7988826815642458
DT Score: 0.8156424581005587
LR Score: 0.8212290502793296
KN Score: 0.8324022346368715
SVC Score: 0.8379888268156425
GNB Score: 0.7821229050279329

```
In [42]: # setup voting classifier and get initial cross val score

vote = VotingClassifier(estimators=[('gb',gb), ('rf',rf), ('et',et), ('ab',ab
), ('dt',dt), ('lr',lr), ('kn',kn), ('svc',svc), ('gnb',gnb)], voting='soft')
vote.fit(features, labels)
cross_val_score(vote, features, labels, cv=5, scoring='accuracy').mean()
```

Out[42]: 0.8283140889719685

In [43]: *# retrain voting classifier and get validation score*

```
vote.fit(features_train, labels_train)
print('Voting Score:', vote.score(features_test, labels_test))
```

Voting Score: 0.8324022346368715

In [44]: *# use itertools combinations to make list of tuples of all possible different combinations of classifiers*

```
clfs = [('gb',gb), ('rf',rf), ('et',et), ('ab',ab), ('dt',dt), ('lr',lr), ('k
n',kn), ('svc',svc), ('gnb',gnb)]
combs = []
```

```
for i in range(2, len(clfs)+1):
    comb = [list(x) for x in itertools.combinations(clfs, i)]
    combs.extend(comb)
```

In [45]: *# Tune voting classifier to use best combination of base classifiers*

```
alg = vote
params = {'estimators': combs}
clf = GridSearchCV(alg, params, cv = 5, scoring = 'accuracy', n_jobs = -1)
clf.fit(features, labels)
print("Best Parameters:", clf.best_params_)
print("Best Score:", clf.best_score_)
vote = clf.best_estimator_
```

```
Best Parameters: {'estimators': [('gb', GradientBoostingClassifier(criterion
='friedman_mse', init=None,
                        learning_rate=0.01, loss='deviance', max_depth=3,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None
e,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_iter_no_change=None, presort='auto',
                        random_state=None, subsample=1.0, tol=0.0001,
                        validation_fraction=0.1, verbose=0,
                        warm_start=False)), ('rf', RandomForestClassifier
(bootstrap=True, class_weight=None, criterion='gini',
      max_depth=None, max_features='auto', max_leaf_nodes=None,
ne,
      min_impurity_decrease=0.0, min_impurity_split=None,
      min_samples_leaf=5, min_samples_split=3,
      min_weight_fraction_leaf=0.0, n_estimators=25,
      n_jobs=None, oob_score=False, random_state=None,
      verbose=0, warm_start=False)), ('dt', DecisionTreeClas
sifier(class_weight=None, criterion='gini', max_depth=None,
      max_features=None, max_leaf_nodes=None,
      min_impurity_decrease=0.0, min_impurity_split=None,
      min_samples_leaf=1, min_samples_split=4,
      min_weight_fraction_leaf=0.0, presort=False,
      random_state=None, splitter='best')), ('kn', KNeighbor
sClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
      metric_params=None, n_jobs=None, n_neighbors=10, p=2,
      weights='uniform')), ('svc', SVC(C=10, cache_size=200, c
lass_weight=None, coef0=0.0,
      decision_function_shape='ovr', degree=3, gamma='auto', kernel='poly',
      max_iter=-1, probability=True, random_state=None, shrinking=True, tol=0.0
01,
      verbose=False))]]}
Best Score: 0.8361391694725028
```

In [46]: *# retrain voting classifier and get validation score*

```
vote.fit(features_train, labels_train)
print('Voting Score:', vote.score(features_test, labels_test))
```

Voting Score: 0.8324022346368715

In [47]: *# retrain voting classifier with full train set, use to make probability predictions and make df of probs, set threshold for probabilities and use to convert probs to predictions*

```
vote.fit(features, labels)
pred_prob = pd.DataFrame(vote.predict_proba(test_df))
threshold = 0.55
y_pred = pred_prob.applymap(lambda x: 1 if x>threshold else 0)
```

In [48]: *# add predictions submission df as survived column, drop all columns but passenger ID and survived, write submission to csv without index to generate submission file*

```
submission['Survived'] = y_pred[1].astype(int)
submission = submission[['PassengerId', 'Survived']]
submission.to_csv('submission.csv', index=False)
```